
Geometrically Consistent Generalizable Splatting

Abstract

Gaussian splatting has emerged as the preferred 3D scene representation due to its incredible speed and accuracy in novel view generation. Various attempts have thus been made in adapting multi-view structure prediction networks to predict per pixel 3D Gaussians from images directly. However, most of the existing work in this domain has focused on enhancing existing depth self-supervised depth prediction networks to estimate additional parameters for 3D Gaussians – orientations, scale, opacity and appearance. This paper shows that view-synthesis loss is insufficient to recover geometrically meaningful splats in this simple manner. We systematically analyse and address the inherent ambiguities in learning 3D Gaussian splats with self-supervision to learn pose-free generalisable splatting. The proposed approach provides state of the art (i) geometrically consistent reconstructions, (ii) relative pose between images and (iii) novel-view synthesis on RealState10K and ACID datasets. We also showcase zero-shot capabilities of the proposed generalizable splatting on ScanNet, where our method substantially outperforms the prior art in recovering geometry and relative pose as well.

1 Introduction

Gaussian splatting [19] has revolutionised 3D structure and appearance modelling from multiple images in the recent past. Departing from traditional depth or point cloud representations for the scene structure 3D Gaussians implicitly model surface reflections and environment lighting to encode view-dependent scene appearance, are memory-efficient compared to explicit volumetric alternatives and facilitate rendering of the scene from arbitrary viewpoints in a fraction of a second, utilising modern GPUs. Owing to the success of 3DGS, a significant effort has been put into improving the 3DGS to work with few-views [48, 4, 14, 32], in the presence of dynamically moving objects [37] and in incorporating object semantics into 3D reconstructions [21]. Real-time simultaneous localisation and mapping approaches have been adapting splats as the inherent scene representation [24, 18]. Splats are also used for the generation of geometrically consistent images and video sequences [34, 37]. 3D Gaussian splats have become the prevalent choice for representing scenes in this era.

The deep learning revolution of the last decade has drastically influenced geometric inference from one or more images. While earlier works in this direction focused on learning neural networks that map image(s) to their depth-maps provided by range sensors [9, 10, 45], recent works have explored reconstructing a registered set of per-pixel point clouds for multiple images to provide state-of-the-art relative pose and structure of the scene [35, 20]. It has also been shown that these feed-forward neural networks to predict geometry from image(s) can be trained without depth sensors, in a self-supervised manner, by minimising view synthesis loss [9, 10, 43, 11]. Structure prediction from single or a few images has been used as the optimisation-free building block in high-fidelity tracking and mapping systems [47, 49]. Learning-assisted recovery of scene geometry has truly become mainstream.

It is natural to marry the incredible capabilities of modern deep neural networks in learning useful geometric priors with the advantages of using 3D Gaussians as the scene representation. Various laudable attempts have been made recently in training neural networks to predict 3D splats directly from images [1, 3, 36, 38, 27, 44, 40]. In the literature, these methods are usually dubbed as Generalizable Gaussian Splatting methods.




Figure 1: **Qualitative comparison of mesh reconstructions from two input views.** We compare the baseline methods MVSplat and NoPoSplat with our approach. The top row displays the textured surface reconstructions; the bottom row shows Gaussian-surface normals for the first input view, along with the rendered RGB image and depth map from a novel (virtual) view. Inaccuracy in rendered depth and normals is evident for both baselines. These inconsistent depths, when fused, create several holes in the mesh reconstruction. Our method recovers accurate geometry of fine details such as the towel, wall painting, and stool.

Most Generalizable Gaussian Splatting methods focus on adapting well-studied structure prediction networks that estimate the means of the Gaussian. Generally, these networks consist of encoders for one or multiple images and decoders to predict Gaussian means. These means are either encoded in the form of per-pixel depths or point clouds, adapting two independent well-studied streams of works for two-view structure prediction via neural networks. Additional decoders are often appended directly to predict the remaining Gaussian parameters (orientation, scale, opacities and view-dependent colour) without much foresight in these methods. All or a partial set of encoders and decoders are learned by usually minimising view-synthesis loss on a few target views by these generalizable splatting methods. This prevalent setup used by Generalizable Gaussian Splatting methods overlooks several key issues that are inherited from the underlying 3DGS optimisation:

- 3D splats are grossly over-parametrised compared to depth-maps or 3D point-clouds. Successful estimation of these 3D splats usually requires a large number of densely sampled viewpoints. 3DGS with few-views is an active field of research, with several solutions relying on heavy regularisation [48, 4]. These developments are ignored in building a successful generalizable Gaussian splatting technique.
- Unlike per-pixel depths or 3D locations, which are uniquely defined for a scene (up to scale), several 3D Gaussian configurations can render input images equally well. Lacking a unique definition, it is unclear how to learn optimal 3D Gaussians even when scene structure is available for supervision in the form of depth maps.
- Successful Gaussian splatting relies on multiple non-differentiable heuristics (such as pruning, splitting, and duplication of Gaussian). Throughout the optimisation, many 3D Gaussians die (reach close to zero opacity), and many new Gaussians are introduced. In the absence of differentiable approximations of these heuristics, existing generalizable splatting approaches only rely on stochastic gradient descent optimisation of a fixed number of Gaussians. i.e. these approaches assume that per-pixel 3D Gaussians will remain perpetually alive during training without any substantial reason.

As a result, the current generalizable Gaussian splatting method often converges to geometrically degenerate scenarios. This is spatially true for Gaussian parameters other than their means – which are relatively stable thanks to adaptation of well well-studied single-view or multi-view depth estimator model initialisations and architectures used for their prediction. Figure 4 showcases that current generalizable splatting frameworks fail to learn meaningful opacities, orientations, or scales of the 3D Gaussians when trained with only self-supervised loss. We can see non-plausible 3D Gaussian

orientations (in the form of normals) as well as unjustifiable elongation of the 3D Gaussian (scales) from both pose-aware and pose-free paradigms of generalizable splatting.

We show that these artefacts are due to over-parametrisation of geometry in the form of splats, which require structural consistency priors to make the self-supervised learning viable. Once these priors are introduced, the Gaussians produced from the proposed method exhibit various geometrically sensible patterns. Briefly, one can see that the proposed method reconstructs a *very low number of splats with small opacities* (near dead Gaussians) compared to baselines, can provide correct surface normals directly from the predicted Gaussian orientations, and produces Gaussian shapes elongated perpendicular to the normal edges and not along the normal edges. We will extensively discuss the geometric inconsistency present in the current Generalizable 3DGS results and how the proposed method alleviates these inconsistencies in section ???. The two approaches we picked for the visualisation of predicted 3D Gaussians broadly represent separate underlying representations for encoding Gaussian means, namely (i) per-pixel depth maps or (ii) per-pixel 3D locations aligned in a single frame of reference. However, to our knowledge, all self-supervised Generalizable 3DGS methods suffer from the same limitations as they inherit representation and loss function from [3, 40].

In this work, we systematically aim to define what the ideal configuration for the geometrically consistent Gaussian should look like and propose appropriate priors to assist Generalized Gaussian Splatting. To that end we opt to work on the recently published NoPo-Splat [40] as our baseline and enhance it. We select No-PoSplat as the baseline because, in addition to its state of performance on multiple geometric tasks, it is a self-supervised approach that does not require ground-truth depth maps, unlike some of the other generalizable splatting adaptations [38]. Additionally, utilising DUST3R [35] framework, the approach is one of the few to provide generalizable splatting from a pair of images without requiring the relative pose of these images.

We address ambiguities in learning over-parametrised 3D Gaussian splats by adding suitable regularisations to the traditionally used view-synthesis loss. Our main observations are: (i) defining the ideal 3D Gaussian orientations to be dominant normals of the scene surfaces (as advocated in [16]) helps resolve structural ambiguities and learning of the Gaussian orientations. We do that by integrating rasteriser from 2DGS framework [16], (ii) when relying on view synthesis loss, ensuring that the 3D Gaussians are pixel-aligned is extremely important. Particularly pose-free methods that use PnP for relative pose estimation between given images require this pixel alignment of Gaussians in ensuring accurate camera pose estimation as well as geometric structure prediction (iii) Jointly learning pixel-aligned normals and 3D location is challenging when self-supervision is used. Often, the optimisation gets stuck in a local minimum with degenerate planar reconstruction despite using existing priors used in [16, 2]. Relying on depth-derived pseudo-normals for orientation supervision provides geometric accuracy.

We trained our generalizable splatting networks on RealState10k (RE10k) dataset [46] adapting pose-free approach of NoPoSplat [40] with appropriate geometric consistency priors. Our method outperform prior art on the tasks of novel-view synthesis, provides a plausible scene structure that can be rendered directly to provide depths from arbitrary views that current methods are incapable of. It also forms a new state of the art in relative pose estimation from two images surpassing existing approaches that include the task-specific methods such as RoMa [6] by a large margin. We also showcase the generalisation capability of our trained models on ScanNet [5] and ACID [23], where our RE10k trained models outperform all existing pose-free generalizable splatting methods in relative-pose estimation and reach the performance of in-domain task-specific pose estimators. Our approach provides the most accurate 3D reconstructions amongst all generalizable splatting approaches in the zero-shot testing setup on Scan-net datasets, including methods that use ground-truth relative pose to infer depth from two views. To the best of our knowledge, it is the first generalizable splatting approach to analyse and evaluate the veracity and meaningfulness of Gaussian Splatting parameters 3D Gaussian orientations and elongations. We believe that the presented analysis pave the way forward for future research training direct 3D Gaussians from images in both depth supervised and self-supervised manner.

2 Related Work

2.1 Self-Supervised Geometry from One or Two Views

Early monocular depth learners minimise photometric error between a view and its reconstruction, jointly estimating depth and ego-motion [9–11, 45, 43]. Multi-view variants build differentiable cost-volumes (MVSNet [39], CAS-MVSNet [12], PatchMatchNet [33]) that outperform classical stereo without laser scans. Recent two-view matchers such as DUST3R [35] and MAST3R [20] fuse image pairs into canonical point clouds; these priors power pose-free approaches like NoPoSplat [40]. Dense SLAM systems (NICE-SLAM [47], NICER-SLAM [49]) further demonstrate the maturity of view-synthesis supervision. Our work adopts this self-supervised paradigm but replaces depth/points with Gaussians, revealing new ambiguities that require additional priors.

2.2 Gaussian Splatting as an Explicit Scene Representation

Per-scene 3D splats. 3DGS [19] introduces anisotropic Gaussians with differentiable rasterisation for real-time view synthesis. Aliasing is alleviated by Mip-Splatting [42]; scene dynamics are handled by 4D-GS [37]. Sparse-view regularisation appears in S2Gaussian [32], while Binocular3DGS [14] leverages stereo consistency. Surface-aligned covariances in SuGaR [13] enable fast mesh extraction, and large-scale models such as LGM [30] push fidelity across datasets.

2D splat variants. 2DGS [16] renders Gaussians in the image plane, linking covariance orientation to surface normals and inspiring our orientation prior.

2.3 Generalisable Gaussian Splatting

Pixel-aligned depth-parametrised models (pose known). PixelSplat [1], MVSplat [3], FreeSplat [36], GPS-Gaussian [44] and Colmap-Free3DGS [8] predict per-pixel Gaussians. HiSplat [31] adds a coarse-to-fine hierarchy, while SplatterImage [29] shows ultra-fast single-view inference. Depth smoothness and redundancy-reduction modules regularise the severe over-parameterisation.

Canonical point-cloud models (weak or no pose). NoPoSplat [40], InstantSplat [7] and Splatt3R [27] regress 3D centres directly in a global frame. eFreeSplat [25] discards epipolar lines, using ViT pre-training and iterative scale alignment. SelfSplat [17] jointly optimises pose and splats in a fully self-supervised loop. Although pose requirements differ, all canonical methods suffer drifting orientations and opacities without extra priors – limitations we explicitly address.

2.4 Regularisation and Geometric Priors

Across both families, authors employ depth smoothness [1], multi-view colour consistency [3], frequency-aware scale constraints [42], opacity sparsity [32], surface-normal alignment [13], and probabilistic sampling to replace non-differentiable pruning [1]. Through dedicated losses for pixel alignment and normal-aware orientation, our framework achieves geometrically consistent generalisable Gaussian splatting.

3 Method

In this section, we introduce our generalizable splitting network, training procedure, and loss function we use to address the ill-posed nature of self-supervised learning in predicting geometrically consistent Gaussian splats from a set of images.

Architecture: Our approach adapts the N-view pose-free generalisable splitting architecture proposed by [40] and consists of three main modules: a transformer-based image encoder, two cross-attention-based asymmetric multi-view feature aggregators and a set of dense prediction transformer decoders to predict parameters of per-pixel Gaussian splats. Specifically, the image encoder transforms RGB image patches into a sequence of tokens and are augmented with the intrinsic parameters of the camera. These pre-image token sequences are processed by a ViT encoder independently. The resulting per-image features are subsequently fused via two sets of feature aggregators, each comprising of multi-head cross-attention layers, effectively aggregating information across multiple

views. These transformer-based networks provide features for the prediction of Gaussian splats for the reference (first) and subsequent images, respectively. Parameters of the Gaussians are estimated by two separate DPT prediction heads: the first head directly infers Gaussian centres for input image pixels, the second head predicts the rest of the Gaussian parameters and combines fine image-level features for accurate appearance estimation. Following [35], we use two asymmetric feature aggregators to provide the Gaussian centres for pixels of each frame to be aligned with the first image of the sequence. As all the recovered Gaussians are aligned, the union of them can be directly rendered to an arbitrary viewpoint without requiring any warping. Notably, both image tokeniser and feature aggregator employ Vision Transformer (ViT) backbones exclusively, without specialised architectures such as epipolar attention or multi-view cost-volumes. This geometry-free architecture does not require relative camera poses during training and inference and has been shown to perform on par or better than the generalisable splitting methods deploying these specialised architectures.

Notations and Problem Definition Assuming we are given with a set of $\mathcal{T}$ context images $\mathcal{I} = \{\mathbf{I}^i \in \mathbb{R}^{H \times W \times 3}\}_{i=1:\mathcal{T}}$, each with known camera intrinsics forming the set $\mathcal{K} = \{\mathbf{K}^i \in \mathbb{R}^{3 \times 3}\}_{i=1:\mathcal{T}}$, capturing a rigid scene. Our objective is to learn a feed-forward neural network with parameters θ that maps these images to a set of 3D Gaussian splats, one splat per pixel for all the input image such that:

$$f_{\theta}(\mathcal{I}, \mathcal{K}) = \left\{ \mathcal{G}_i^j := (\boldsymbol{\mu}_j^i, \alpha_j^i, \mathbf{q}_j^i, \mathbf{s}_j^i, \mathbf{c}_j^i) \right\}_{i=1:\mathcal{T}, j=1:H \times W} \quad (1)$$

where $\mathcal{G}_i^j$ is the 3D Gaussian corresponding to a pixel i in image j , characterized by its: 3D center $\boldsymbol{\mu}_j^i \in \mathbb{R}^3$; orientation represented by a unit quaternion vector $\mathbf{q} \in \mathbb{R}^4$; and two scale parameters $\mathbf{s} \in \mathbb{R}^2$ defining the elongation of 3D Gaussian splat; opacity $\alpha \in \mathbb{R}$; and appearance encoded in the form of coefficients of the spherical harmonic function of order d , $\mathbf{c} \in \mathbb{R}^d$. In this work, we advocate the use of per-pixel 2D Gaussians to represent the scene instead of the standard 3D Gaussian splats. Following [16], we assume that estimated Gaussians are aligned with the scene surface and its elongation perpendicular to the local surface normal is zero. We show through extensive evaluations how this choice help Generalizable Gaussian splitting.

Both, Gaussian centre $\boldsymbol{\mu}_j^i$ and orientations $\mathbf{q}_j^i$ are defined in the image coordinates of the first image from the set. Given these predicted $M \times N \times \mathcal{T}$ Gaussians¹, novel views $C^f \forall f \in \{1, \dots, \mathcal{F}\}$ of the scene can be rendered from viewpoints defined by the projection matrix $P^f = \mathbf{K}^f[\mathbf{R}^f|\mathbf{T}^f]$ to be matched with observed images I^f s. Like [40], we propose to minimise the view synthesis loss from the predicted 2D Gaussian splats:

$$L_{synthesis} = \sum_{i=1}^{\mathcal{F}} \mathcal{L}(I^f, C^f); \text{ where, } C^f(u, v) = \sum_{k=1}^K c_k \alpha_k \mathcal{G}_{2D}(u, v) \prod_{j=1}^{k-1} (1 - \alpha_j \mathcal{G}_{2D}(u, v)) \quad (2)$$

where $\mathcal{L}(a, b)$ represents the absolute sum of colour differences between observed and rendered images combined with the LPIPS error. $C^f(u, v)$ is the color of the pixel u, v in image C^f rendered by blending K ordered 2D Gaussians using the 2DGS rasterizer. $\mathcal{G}_{2D}$ is the projection of the 2D Gaussians covariance matrix Σ_{2D} on the image plane of image i (See [16] for details).

As shown in section 5, this view synthesis loss is proven to be insufficient for learning geometrically meaningful splats. We propose to minimise two additional regularisation losses: (i) Depth-surface normal consistency term $\mathcal{L}_{normal}$ aligns 2D splat orientations with the rendered depth, and a grid alignment loss $\mathcal{L}_{grid}$ to ensure that the estimated splats are aligned with the pixels of the provided images. These two losses, alongside the view synthesis loss, form our training objective:

$$\mathcal{L} = \mathcal{L}_{synthesis} + \lambda_o \mathcal{L}_{orient} + \lambda_a \mathcal{L}_{align}, \quad (3)$$

where λ_o and λ_a are weighting factors balancing the influence of each regularization. The motivation, formulation and impact of the regularisation term have been discussed in the following sections.

3.1 Learning Gaussian Orientations.

Both pose-free and pose-aware generalised Gaussian splitting approaches fail in learning meaningful orientations of the splats. To provide geometric meaning for the splat orientations, we propose

¹Like most other methods for generalizable splitting, one can predict more than one splat per pixel in our framework as well but we restrict our experiments to one splat per pixel only.

aligning them with the dominant surface normals of the scene they belong to. To that end, we follow the setup in [16] to estimate only two non-zero splat scales– instead of the usual three – and enforce the elongation of splats in the direction of surface normal to be zero. In particular, we model the rank-deficient Gaussian covariance matrix with its null space corresponding to the surface normal vector.

Swapping 3D splats to 2D surface elements reduces the over-parameterisation to an extent; however, the view synthesis loss of eq. (2) does not give a strong enough supervision signal for orientations. The most natural way to supervise Gaussian orientation is to use surface normal regularisation from [16]. Authors propose to render 2D splat orientations in the form of a normal map with alpha blending and force these normals to be consistent with the rendered depth for a training view. This consistency term aligns 2D splats that contribute to the rendering of a single pixel to be perpendicular to the pixel’s rendered normal. While such prior has worked well in [16] and its variants used for few-shot Gaussian splitting techniques, we observe that it can not be used naively for generalisable splitting. Adding this regularisation for joint splat position and orientation estimation locks the generalizable splitting network into a local minimum, producing near-planar scenes. Detaching the rendered depth from the computation graph helps fix this problem. This setup can be seen as using depth rendering during optimisation to generate pseudo-labels for orientations of 2D splats and produce reasonable splat orientation and locations. It should be noted that successful training of the Gaussian orientation with such normal-loss still requires L_{grid} . The results of the [16] inspired loss for learning orientation are discussed in the supplementary material.

However, we opt to use a rather simple alternative to supervise the splat orientations in our approach. Assuming that the estimated splats $\mathcal{S}_i^j$ are aligned with the 2D image pixel $j = (u, v)$ in image i , we derive the local surface normal for this splat using the relative 3D locations of its local neighbourhood in image i . We propose to minimise the following normal loss:

$$L_{orient} = \frac{1}{\mathcal{T}(M-2)(N-2)} \sum_{i=1}^{\mathcal{T}} \sum_{u=2}^{M-1} \sum_{v=2}^{N-1} \|1 - \langle N_{u,v}^i, \hat{N}_{u,v}^i \rangle\|_{\epsilon} \quad (4)$$

$$\hat{N}_{u,v}^i = \|(\mu_{(u+1,v)}^i - \mu_{(u-1,v)}^i) \times (\mu_{(u,v+1)}^i - \mu_{(u,v-1)}^i)\|_* \quad (5)$$

where, $N_{u,v}^i$ is the null space of the Gaussian covariance matrix defining the normal to the plane defined by 2D splat corresponding to pixel (u, v) in image i , $\mu_{(u,v)}^i$ is the estimated 3D location of this splat’s centre, $\langle \cdot, \cdot \rangle$ is the dot product of two vectors, $\|\cdot\|_{\epsilon}$ is Huber loss (implemented as SmoothL1Loss in pytorch) and $\|\cdot\|_*$ represents vector normalization. Here we ignore the normal computation at the corner of the image and use central differences in computing local structure derivatives, however alternative padding and normal approximations can be used.

Note that, unlike the loss used in 2DGS [16] which involves rasterization of normals as well as depth to enforce consistency, the L_{orient} here does not involve any rasterisation. It is motivated by the simple observation that while depth supervised frameworks like DUST3R [35] predict registered 3D point clouds corresponding to the pixels of input images, these point clouds are directly supervised from back-projected (and warped) depthmaps. By extension, if the ground truth depths were availed for learning 3D/2D splats, the orientations can be directly supervised by the (warped) normal maps as well. This simple but effective way to directly supervised surface normals provides best generalizable splitting results under self-supervision. However, the same supervision can be used to improve orientation estimation in depth supervised splat learning frameworks such as [38] as well.

3.2 Pixel aligned Gaussian splats:

Although the first generalizable splitting approach worked in a pose-aware setup and adapted two-view depth predictor architectures, they by construction were designed to provide each 3D Gaussians to lie on a specific viewing ray. Pose free approaches aimed at directly estimating the 3D Gaussians locations as output of a DPT decoder head. While this alleviated the need for warping 3D Gaussian splats with known cameras, the parametrization rendered the structure estimation problem ill-posed especially under self supervised setup. Specifically, unlike depth supervision used in DUST3R [35] that learns a structural bias that the 3D point cloud projects to regular grid in the input image, view synthesis loss does not provide any such bias. 3D Gaussians thus can move freely to geometrically

degenerate configurations hampering both structure and relative pose estimation of pose free methods. We enforce that the 3D Gaussians corresponding to a pixel of the input frame should lie on the corresponding viewing ray. To that end, we project centres $\mu_{u,v}^i$'s of reconstructed 3D Gaussians with known camera location and intrinsic matrix to match the location of pixels forming regular grids. We minimize the following loss:

$$L_{align} = \frac{1}{\sum_{i,u,v} \mathcal{M}_{u,v}^i} \sum_{i=1}^T \sum_{u=1}^M \sum_{v=1}^N \mathcal{M}_{u,v}^i \| [u, v]^T - \Pi(K^i [R^i | T^i] \mu_{u,v}^i) \|_2^2, \quad (6)$$

where $\mathcal{M}_{u,v}^i$ is an indicator function defining splats getting projected inside the image and having positive depths and $([X, Y, Z]^T) = [X/Z, Y/Z]^T$ is the projection function. The proposed loss plays a crucial role in Pnp-based relative pose estimation as well as accurate structure estimation.

4 Analysis of figure 1

It is evident in Figure 4 that none of the existing Generalizable 3DGS approaches provide meaningful orientations or scales for the estimated Gaussians. In particular, we visualize the normal of the dominant plane of 3D Gaussians that corresponds to the direction in which the Gaussians is minimally elongated. This minimal elongation is labelled as scale 3 in the figure and quantifies the uncertainty in localising the scene surface around this Gaussian. The estimated normals from the existing Generalizable 3DGS methods lack any geometric meaning whereas the proposed method can provide meaningful normal maps conforming to plausible scene geometry. For example, notice multiple planes confirming to Manhattan prior in the scene in Figure 4 and the corresponding normal predictions.

The existing Generalised 3DGS methods does not provide geometrically meaningful estimates of the shape (scales) to the Gaussians either. One can observe clear over-fitting to view synthesis loss – which is highly sensitive to the location of image edges – in apparent image texture in the scales. Higher values in scale 3 suggest billboard-like elongated Gaussians perpendicular to the scene surface that existing Generalizable 3DGS approaches produce. On the other hand, our approach uses 2DGS rasterizer by construction, forcing this scale to be 0, eliminating these billboard-like Gaussians. Not only does the proposed approach localize scene surfaces, but the dominant structural edges are visible in the estimated scales. For example, look at the interesting case of the edge of the wall at the intersection of two corridors in Figure 4. The Gaussians estimated by our method are very skewed in one direction. This is expected out of geometrically consistent Gaussians to have a very narrow extent perpendicular to geometric edges, whereas an arbitrarily large extent along the edge. It can also be seen that more planar areas have more homogeneously elongated Gaussians. No clear geometric edge or homogeneity in the size and shape of Gaussians can be found in the existing self-supervised approaches.

We can also see the opacities of the reconstructed Gaussians. Large regions of the image have very low opacities, pointing to many near-dead Gaussians estimated by the existing Generalizable 3dgs methods. These low opacity Gaussians are often found in the regions of shadows, light sources and reflection, suggesting that instead of modelling these variations with spherical harmonics as one would expect, the view-dependent colour is been modelled by changing opacities. In contrast, it can be seen that the proposed method consists of less number of near-dead Gaussians and low-opacity regions corresponding to image edges and shadows.

5 Experiments

5.1 Experimental Setup and Details

Datasets: To evaluate geometry and novel-view synthesis performance, we follow common practices in generalizable splatting methods [1, 3, 40] by training our model on the large-scale RealEstate10K [46] (RE10k) dataset. We use standard data splits on RE10K dataset and evaluate novel-view synthesis and pose estimation performance on the dataset for quantitative comparisons. ACID [23] and ScanNet [5] are used for evaluating the generalisation capability of the our trained network on both relative pose and geometry estimation, where applicable.

RE10k comprises real estate videos (predominantly indoor scenes) downloaded from YouTube, and is divided into 67,477 training scenes and 7,289 testing scenes following the same train-test

Table 1: **Pose estimation performance (AUC) at multiple error thresholds on RE10k [46] (in-domain) and on ScanNet-V1 [5] and ACID [23] (cross-domain).** Our method delivers state-of-the-art results across all thresholds and datasets—even though it is trained *only* on RE10k. On ScanNet, in particular, it surpasses the baseline while relying solely on PnP+RANSAC, whereas the baseline requires additional photometric pose refinement. The overall best results are highlighted in **bold**, and the best PnP+RANSAC-only result is underlined. Note that approaches marked † are trained on additional data (including ScanNet and ACID) and marked ‡ use additional supervision such as ground truth depth.

Method	RE10k			ScanNet-V1			ACID		
	5° ↑	10° ↑	20° ↑	5° ↑	10° ↑	20° ↑	5° ↑	10° ↑	20° ↑
CoPoNeRF†	0.161	0.362	0.575	-	-	-	0.078	0.216	0.398
DUS3R†	0.301	0.495	0.657	-	-	-	0.166	0.304	0.437
MASt3R†	0.372	0.561	0.709	-	-	-	0.234	0.396	0.541
RoMa‡	0.546	0.698	0.797	0.168	0.361	0.575	0.463	0.588	0.689
<i>PnP+RANSAC only</i>	NoPoSplat	0.572	0.728	0.833	0.078	0.198	0.394	0.337	0.497
	Ours (2DGS)	0.588	0.737	0.832	0.085	0.223	0.432	0.344	0.513
	Ours (2DGS+Grid)	0.621	0.760	0.849	0.123	0.279	0.471	0.382	0.540
	Ours (2DGS+Normal)	0.613	0.756	0.848	0.118	0.267	0.460	0.376	0.537
	Ours (2DGS+Grid+Normal)	<u>0.627</u>	<u>0.766</u>	<u>0.855</u>	<u>0.135</u>	<u>0.289</u>	<u>0.479</u>	<u>0.392</u>	<u>0.547</u>
<i>w/ Photometric Optimisation</i>	NoPoSplat	0.672	0.791	0.868	0.109	0.256	0.463	0.456	0.593
	Ours (2DGS)	0.672	0.788	0.859	0.129	0.298	0.515	0.460	0.599
	Ours (2DGS+Grid)	0.686	0.799	0.870	0.136	0.311	0.512	0.474	0.607
	Ours (2DGS+Normal)	0.679	0.798	0.871	0.141	0.323	0.520	0.475	0.610
	Ours (2DGS+Grid+Normal)	0.689	0.804	0.876	0.156	0.334	0.539	0.488	0.619

Table 2: **Depth estimation performance for novel views on ScanNet-V1 [5].** Two input camera views are used to reconstruct the scene using Gaussians. We propose using the accuracy of rendered depth for novel views as a holistic measure of the model’s interpolation capability for 3D reconstruction. Our method outperforms all other approaches across all evaluated metrics by a clear margin. The best overall results are highlighted in **bold**, and the best-performing method without camera pose refinement is indicated by underline.

Method	Rendered Depth (Novel Views)		
	Abs Rel↓	$\delta_1 < 1.10$ ↑	$\delta_1 < 1.25$ ↑
pixelSplat	0.299	0.552	0.818
MVSplat	0.189	0.412	0.745
<i>w/o Pose Refine.</i>	NoPoSplat	0.131	0.554
	NoPoSplat*	0.135	0.552
	Ours (2DGS)	0.121	0.668
	Ours (2DGS+Grid)	0.109	0.679
	Ours (2DGS+Normal)	0.115	0.674
	Ours (2DGS+Grid+Normal)	<u>0.108</u>	<u>0.680</u>
<i>w Pose Refine.</i>	NoPoSplat	0.126	0.567
	NoPoSplat*	0.130	0.564
	Ours (2DGS)	0.114	0.692
	Ours (2DGS+Grid)	0.102	0.706
	Ours (2DGS+Normal)	0.106	0.704
	Ours (2DGS+Grid+Normal)	0.100	0.707

split as in previous works [1, 3]. ACID contains mainly aerial nature scenes captured by drones. For both datasets, camera poses were computed using COLMAP [26]. In contrast, ScanNet [5] consists of RGB-D captures of indoor scenes with different camera motions and scene characteristics compared to the other two datasets. We use depths and relative camera poses available in two test splits: ScanNet-V1 – comprising of 2000 image pairs for generalised splitting and ScanNet-1500 – comprising 1500 test image pairs.




Figure 2: **Qualitative comparison of rendered novel-view depth on RE10k [46] (top two rows) and ACID [23] (bottom row).**




Figure 3: **Qualitative comparison of rendered novel-view depth on ScanNet-V1 [5].**

Implementation details. All of our generalizable splat predictors are trained using the setup described in NoPoSplat [40]. Training is performed on a cluster of 24 A100 (40 GB) GPUs with a per-GPU batch size of 6 (total batch size 144), while all evaluations in Section 5 are run on a single A6000 GPU. Models are trained on the RE10k training split for 18751 iterations using 256×256 images and four rendered virtual views to minimise the view-synthesis loss. Regularisation weights are set to $\lambda_o = 0.05$ and $\lambda_a = 0.1$ for the full model and all ablations, where applicable. We use a base learning rate of 2×10^{-4} , while layers of the DUS3R backbone are updated with a reduced rate of 2×10^{-5} . We also note that, like NoPoSplat, DUS3R-based splitting extensions are unable to be




Figure 4: Qualitative comparison of Gaussians. The first row visualises our 2D Gaussians with the second and third containing results from NoPoSplat and MVSplat, respectively. All visualised Gaussians correspond to the context image (a). Row 1: (b-c) shows the extent of 2D Gaussians along the estimated scene surface it belongs to. The surface is defined by its normal, which is shown in (d). Row 2/3 for NoPoSplat and MVSplat respectively: (f/k) show minimal Gaussian elongation. This direction is assumed to be perpendicular to the dominant surface and the elongation quantises the uncertainty in the surface localisation. (g/l, h/m) show Gaussian elongation along the estimated dominant surface, i.e. comparable to corresponding columns in top row, (i/n) in the two rows shows estimated normals of the dominant surface for baseline methods. Additionally, (e,j,o) show the estimated opacity of the Gaussians.

Table 3: **Novel-view synthesis performance on the RealEstate10k dataset [46].** We compare pose-required and pose-free methods against our proposed variants. “NoPoSplat” refers to results from publicly available checkpoints, while “NoPoSplat*” indicates results from our retrained baseline using identical hyperparameters for fair comparison. The best results among pose-free methods (and without optimizing camera poses for target views) are highlighted in **bold**, and the best-performing pose-required method is indicated by underline.

		Small			Medium			Large			Average		
Method		PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓
Pose-Required	pixelNeRF	18.417	0.601	0.526	19.930	0.632	0.480	20.869	0.639	0.458	19.824	0.626	0.485
	AttnRend	19.151	0.663	0.368	22.532	0.763	0.269	25.897	0.845	0.186	22.664	0.762	0.269
	pixelSplat	20.263	0.717	0.266	23.711	0.809	0.181	27.151	0.879	0.122	23.848	0.806	0.185
	MVSplat	<u>20.353</u>	<u>0.724</u>	<u>0.250</u>	<u>23.778</u>	<u>0.812</u>	<u>0.173</u>	<u>27.408</u>	<u>0.884</u>	<u>0.116</u>	<u>23.977</u>	<u>0.811</u>	<u>0.176</u>
Pose-Free	DUS3R	14.101	0.432	0.468	15.419	0.451	0.432	16.427	0.453	0.402	15.382	0.447	0.432
	MASt3R	13.534	0.407	0.494	14.945	0.436	0.451	16.028	0.444	0.418	14.907	0.431	0.452
	Splatt3R	14.352	0.475	0.472	15.529	0.502	0.425	15.817	0.483	0.421	15.318	0.490	0.436
	CoPoNeRF	17.393	0.585	0.462	18.813	0.616	0.392	20.464	0.652	0.318	18.938	0.619	0.388
	NoPoSplat	21.086	0.721	0.237	23.134	0.776	0.185	25.086	0.818	0.141	23.189	0.775	0.185
	NoPoSplat*	21.097	0.723	0.237	23.191	0.779	0.187	25.107	0.817	0.144	23.244	0.778	0.187
	Ours (2DGS)	21.051	0.725	0.245	23.338	0.784	0.187	25.541	0.830	0.140	23.407	0.782	0.188
	Ours (2DGS+Grid)	21.341	0.736	0.236	23.453	0.788	0.185	25.529	0.829	0.141	23.528	0.787	0.185
	Ours (2DGS+Normal)	21.275	0.733	0.236	23.198	0.780	0.187	25.205	0.823	0.144	23.300	0.781	0.187
	Ours (2DGS+Grid+Normal)	21.349	0.739	0.234	23.463	0.789	0.184	25.628	0.832	0.141	23.564	0.789	0.184

trained from scratch on RE10k without distillation of the weights; instead, we initialise them with MASt3R-pretrained weights, and allow other baselines to use similar supervised initialisation of the backbone where available. For a fair comparison with the baseline, we report two versions of NoPoSplat: the publicly released checkpoint (“NoPoSplat”) and a retrained variant (“NoPoSplat*”) obtained with the *exact* hyper-parameters listed above. The source code and pre-trained models are publicly available at <https://github.com/m80hz/Gen2DGS>.

5.2 Results and Analysis

We evaluate novel-view synthesis, the relative pose between the input images and the reconstructed geometry, where applicable, to benchmark our generalizable splatting network. Ours is a pose-free approach where two images alongside the intrinsic camera parameters are used to generate splats. The camera coordinate frame for the first image is used as the de-facto coordinate frame.

5.2.1 Relative pose estimation

Relative pose is evaluated by computing the area under the cumulative pose error curve with three different thresholds. We report results deploying a simple PnP + RANSAC algorithm to align the splat means with image grid as proposed in DUS3R[35]. NoPosplat [40] proposes a gradient descent-based relative pose refinement scheme in which the predicted Gaussian Splats are rendered to generate optimal input image pairs for refining the pose obtained by PnP+RANSAC. Pose Jacobians from [?] are used for this refinement for a fixed number of iterations. Although these pose Jacobians are not readily usable with 2D Gaussian splats, we convert our estimated 2D Gaussians to their 3D counterparts by augmenting a small non-zero scale to facilitate a comparison with the relative pose estimation results reported in [40] as well.

Table 1 reports the quantitative comparison of relative pose estimation for several methods. CoPoNeRF [15] is trained on RE10k and ACID with explicit supervision on camera poses; DUS3R [35] relies on representative indoor RGB-D and Internet SfM imagery (e.g., ScanNet++ [41] and MegaDepth [22]) and employs a 3D regression loss that directly supervises both depth and pose; MASt3R [20] follows the same supervision scheme to DUS3R but broadens the data mix with large-scale outdoor sequences such as Waymo [28]; RoMA [6], trained on MegaDepth and ScanNet, likewise uses depth-and-pose supervision. Our models, on the contrary, does not use explicit depth supervision and are only trained on Re10k data. Despite a lack of depth supervision and smaller training dataset, the trained models outperform these methods by a large margin – both on in domain testset of Re10k and on ACID and Scannet++. The only exception is RoMa’s performance on ScanNet-V1 – the dataset on which the method was trained for relative pose estimation explicitly.

In particular, when compared to baseline Nopo-Splat, our method provides significant improvements in relative pose estimation when a simple PnP+RANSAC algorithm is used. In particular, the alignment loss L_{align} made a significant difference, both in-domain and out-of-domain zero-shot testing. The splat orientation loss L_{orient} also helped pose estimation a bit, with our full loss giving a further boost to performance. As reported in NoPoSplat, the input image synthesis loss minimisation improved the pose estimation for our models too. While the improvement on the Re10K and ACID dataset was limited in our experiments, we see a proportionate improvement on the Scannet-V1 data with the optimisation.

5.2.2 Evaluating the geometry

Geometric veracity of the estimated 2D/3D Gaussian splats is the key focus of this work. Traditionally, the geometry predicted by feed-forward neural networks is evaluated by measuring the depth errors for the input views. However, input depths do not capture the interpolation capability of splats, and are insensitive to the opacity, orientation, and scale parameters. We propose a more holistic evaluation of the scene structure by rendering multiple virtual depth maps from the reconstructed splats and reporting the average Absolute Relative Error and Accuracy. As we do not focus on extrapolating the scene beyond the given view frustum, we keep use the view-synthesis evaluation testset for depth evaluation as well. These virtual depths are reconstructed using the ground-truth relative pose with respect to the first input frame in our experiments, assuming that the Multi-view Gaussians are perfectly aligned. This puts pose-free methods at a severe disadvantage, as small pose alignment errors can result in large depth errors. However, despite this disadvantage, pose-free methods outperformed pose-aware counterparts by a large margin, as shown in Table 2. The proposed approach additionally outperforms the NoPoSplat baseline by a large margin as well. While 2DGS parametrisation helped in accurate structure recovery, larger gains were made with orientation and alignment losses. The same trend persisted when per-image pose optimisation was used to align the recovered splats with the target images while improving the performance for all the models.

A more accurate evaluation of reconstruction can be devised by rendering a large number of virtual depth maps from the generated Gaussian splats and fusing them back to create scene’s truncated signed distance field. 3D surface from the TSDF can be extracted using the marching cube algorithm to be compared with the ground truth, where available. We use such virtual-depth-rendering $\rightarrow$ TSDF fusion $\rightarrow$ surface extraction for visual comparison facilitated in the figure 1. An appropriate distance measure, such as Chamfer distance between ground truth and generated mesh, is used as a good measure for pose error normalised surface reconstruction accuracy in the supplementary material.

As the Re10k and ACID datasets do not have depth ground truth, we rely on visual comparison. Figure 2 visualises the in-domain and out-of-domain depths for pixelSplat, MVSplat and NoPoSplat on rendering novel view depthmaps. It can be seen that MVSplat and NoPoSplat depths were hypersensitive to image texture, while pixelSplat provided much noisier depth maps. Our method produces more plausible depths despite not demanding the relative pose between images. Similar depth errors are observed in Figure 3 when we test these generalisable splitting methods on the ScanNet test instances.

While the focus of this work in learning generalizable splats, we report comparisons with prior art for depth estimation networks in supplementary material. As these methods cannot render virtual view depths used to evaluate the reconstruction, we propose to pose-align our recovered splats with individual input images by minimising the view synthesis loss.

5.2.3 Novel view synthesis

While not germane to the proposed work, we compare in-domain and out-of-domain novel view synthesis performance with the relevant baseline methods in Tables 3 and ??, respectively. Novel-views are synthesised by rendering all reconstructed splats for pose-free methods, given the fixed target pose w.r.t the first input image. For methods that estimate splat locations in the form of depth maps, an additional warping step with ground-truth relative camera pose is used to generate a set of splats aligned with the coordinate frame of the first image. Note that we do not opt for optimising camera pose for optimal target view synthesis. While many NeRF and splats method reports view synthesis errors by target image pose optimisation, it is well known that geometric inconsistencies can be hidden by such pose optimisation easily. We opt not to peek into the answers in synthesising novel views and report results without pose optimisation. We report PSNR, SSIM and LPIPS to

quantify view synthesis performance. It can be seen that the proposed method outperforms prior art in novel-view synthesis largely due to the warping-free nature of the proposed method. A small gain over Noposplat on novel-view synthesis was seen when the proposed loss was used for training. It is noted that the grid loss does not improve out-of-domain novel view performance and at times hampers synthesis accuracy when used in isolation – especially when tested on images out of domain (e.g. scannet).

6 Conclusion

We present a novel self supervise generalisable splitting network that mitigate several geometric inconsistencies in splat recovery which were overlooked by the community. Our network provides state of the art, geometrically consistent splats from two unposed images. While we trained our network on Re10k, using asymmetric transformer based architecture under self-supervision, the key contributions of the work are invariant to these choices. The presented priors in the work will help future development on generalisable splitting and learning-based 3D scene recovery.

References

- [1] David Charatan, Sizhe Lester Li, Andrea Tagliasacchi, and Vincent Sitzmann. pixelsplat: 3d gaussian splats from image pairs for scalable generalizable 3d reconstruction. In *CVPR*, 2024. 1, 4, 7, 8
- [2] Danpeng Chen, Hai Li, Weicai Ye, Yifan Wang, Weijian Xie, Shangjin Zhai, Nan Wang, Haomin Liu, Hujun Bao, and Guofeng Zhang. Pgsr: Planar-based gaussian splatting for efficient and high-fidelity surface reconstruction. *TVCG*, 2024. 3
- [3] Yuedong Chen, Haofei Xu, Chuanxia Zheng, Bohan Zhuang, Marc Pollefeys, Andreas Geiger, Tat-Jen Cham, and Jianfei Cai. Mvsplat: Efficient 3d gaussian splatting from sparse multi-view images. In *ECCV*, 2024. 1, 3, 4, 7, 8
- [4] Jaeyoung Chung, Jeongtaek Oh, and Kyoung Mu Lee. Depth-regularized optimization for 3d gaussian splatting in few-shot images. In *CVPR*, 2024. 1, 2
- [5] Angela Dai, Angel X Chang, Manolis Savva, Maciej Halber, Thomas Funkhouser, and Matthias Nießner. Scannet: Richly-annotated 3d reconstructions of indoor scenes. In *CVPR*, 2017. 3, 7, 8, 9
- [6] Johan Edstedt, Qiyu Sun, Georg Bökman, Mårten Wadenbäck, and Michael Felsberg. Roma: Robust dense feature matching. In *CVPR*, 2024. 3, 11
- [7] Zhiwen Fan, Wenyan Cong, Kairun Wen, Kevin Wang, Jian Zhang, Xinghao Ding, Danfei Xu, Boris Ivanovic, Marco Pavone, Georgios Pavlakos, et al. Instantsplat: Unbounded sparse-view pose-free gaussian splatting in 40 seconds. *arXiv preprint arXiv:2403.20309*, 2024. 4
- [8] Yang Fu, Sifei Liu, Amey Kulkarni, Jan Kautz, Alexei A Efros, and Xiaolong Wang. Colmap-free 3d gaussian splatting. In *CVPR*, 2024. 4
- [9] Ravi Garg, Vijay Kumar Bg, Gustavo Carneiro, and Ian Reid. Unsupervised cnn for single view depth estimation: Geometry to the rescue. In *ECCV*, 2016. 1, 4
- [10] Clément Godard, Oisín Mac Aodha, and Gabriel J Brostow. Unsupervised monocular depth estimation with left-right consistency. In *CVPR*, 2017. 1
- [11] Clement Godard, Oisín Mac Aodha, Michael Firman, and Gabriel J. Brostow. Digging into self-supervised monocular depth estimation. In *ICCV*, 2019. 1, 4
- [12] Xiaodong Gu, Zhiwen Fan, Siyu Zhu, Zuozhuo Dai, Feitong Tan, and Ping Tan. Cascade cost volume for high-resolution multi-view stereo and stereo matching. In *CVPR*, 2020. 4
- [13] Antoine Guédon and Vincent Lepetit. Sugar: Surface-aligned gaussian splatting for efficient 3d mesh reconstruction and high-quality mesh rendering. *CVPR*, 2024. 4
- [14] Liang Han, Junsheng Zhou, Yu-Shen Liu, and Zhizhong Han. Binocular-guided 3d gaussian splatting with view consistency for sparse view synthesis. *NeurIPS*, 2024. 1, 4
- [15] Sunghwan Hong, Jaewoo Jung, Heeseong Shin, Jiaolong Yang, Seungryong Kim, and Chong Luo. Unifying correspondence pose and nerf for generalized pose-free novel view synthesis. In *CVPR*, 2024. 11

- [16] Binbin Huang, Zehao Yu, Anpei Chen, Andreas Geiger, and Shenghua Gao. 2d gaussian splatting for geometrically accurate radiance fields. In *ACM SIGGRAPH 2024*, 2024. 3, 4, 5, 6
- [17] Seunghyun Kang, Hyunwoo Lee, and Hyeongju Chae. Selfsplat: Pose-free and 3d-prior-free generalizable 3d gaussian splatting. *CVPR*, 2025. 4
- [18] Nikhil Keetha, Jay Karhade, Krishna Murthy Jatavallabhula, Gengshan Yang, Sebastian Scherer, Deva Ramanan, and Jonathon Luiten. Splatam: Splat, track & map 3d gaussians for dense rgb-d slam. In *CVPR*, 2024. 1
- [19] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM TOG*, 2023. 1, 4
- [20] Vincent Leroy, Yohann Cabon, and Jérôme Revaud. Grounding image matching in 3d with mast3r. *arXiv preprint arXiv:2406.09756*, 2024. 1, 4, 11
- [21] Mingrui Li, Shuhong Liu, Heng Zhou, Guohao Zhu, Na Cheng, Tianchen Deng, and Hongyu Wang. Sgs-slam: Semantic gaussian splatting for neural dense slam. In *ECCV*, 2024. 1
- [22] Zhengqi Li and Noah Snavely. Megadepth: Learning single-view depth prediction from internet photos. In *CVPR*, 2018. 11
- [23] Andrew Liu, Richard Tucker, Varun Jampani, Ameesh Makadia, Noah Snavely, and Angjoo Kanazawa. Infinite nature: Perpetual view generation of natural scenes from a single image. In *ICCV*, 2021. 3, 7, 8, 9
- [24] Hidenobu Matsuki, Riku Murai, Paul HJ Kelly, and Andrew J Davison. Gaussian splatting slam. In *CVPR*, 2024. 1
- [25] Zhiyuan Min, Yawei Luo, Jianwen Sun, and Yi Yang. Epipolar-free 3d gaussian splatting for generalizable novel view synthesis. *NeurIPS*, 2024. 4
- [26] Johannes L Schonberger and Jan-Michael Frahm. Structure-from-motion revisited. In *CVPR*, 2016. 8
- [27] Brandon Smart, Chuanxia Zheng, Iro Laina, and Victor Adrian Prisacariu. Splatt3r: Zero-shot gaussian splatting from uncalibrated image pairs. *arXiv preprint arXiv:2408.13912*, 2024. 1, 4
- [28] Pei Sun, Henrik Kretschmar, Xerxes Dotiwalla, Aurelien Chouard, Vijaysai Patnaik, Paul Tsui, James Guo, Yin Zhou, Yuning Chai, Benjamin Caine, et al. Scalability in perception for autonomous driving: Waymo open dataset. In *CVPR*, 2020. 11
- [29] Stanislaw Szymanowicz, Christian Rupprecht, and Andrea Vedaldi. Splatter image: Ultra-fast single-view 3d reconstruction. In *CVPR*, 2024. 4
- [30] Jiaxiang Tang, Zhaoxi Chen, Xiaokang Chen, Tengfei Wang, Gang Zeng, and Ziwei Liu. Lgm: Large multi-view gaussian model for high-resolution 3d content creation. *arXiv preprint arXiv:2402.05054*, 2024. 4
- [31] Shengji Tang, Weicai Ye, Peng Ye, Weihao Lin, Yang Zhou, Tao Chen, and Wanli Ouyang. Hisplat: Hierarchical 3d gaussian splatting for generalizable sparse-view reconstruction. *ICLR*, 2025. 4
- [32] Ziyu Wan, Hao Gao, Rui Xiong, and Fang Du. S2gaussian: Sparse-view super-resolution 3d gaussian splatting. In *CVPR*, 2025. 1, 4
- [33] Fangjinhua Wang, Silvano Galliani, Christoph Vogel, Pablo Speciale, and Marc Pollefeys. Patchmatchnet: Learned multi-view patchmatch stereo. In *CVPR*, 2021. 4
- [34] Qianqian Wang, Yifei Zhang, Aleksander Holynski, Alexei A. Efros, and Angjoo Kanazawa. Continuous 3d perception model with persistent state. In *CVPR*, 2025. 1
- [35] Shuzhe Wang, Vincent Leroy, Yohann Cabon, Boris Chidlovskii, and Jerome Revaud. Dust3r: Geometric 3d vision made easy. In *CVPR*, 2024. 1, 3, 4, 5, 6, 11
- [36] Yunsong Wang, Tianxin Huang, Hanlin Chen, and Gim Hee Lee. Freesplat: Generalizable 3d gaussian splatting towards free view synthesis of indoor scenes. *NeurIPS*, 2024. 1, 4
- [37] Guanjun Wu, Taoran Yi, Jiemin Fang, Lingxi Xie, Xiaopeng Zhang, Wei Wei, Wenyu Liu, Qi Tian, and Xinggang Wang. 4d gaussian splatting for real-time dynamic scene rendering. In *CVPR*, 2024. 1, 4
- [38] Haofei Xu, Songyou Peng, Fangjinhua Wang, Hermann Blum, Daniel Barath, Andreas Geiger, and Marc Pollefeys. Depthsplat: Connecting gaussian splatting and depth. In *CVPR*, 2025. 1, 3, 6

- [39] Yao Yao, Zixin Luo, Shiwei Li, Tian Fang, and Long Quan. Mvsnet: Depth inference for unstructured multi-view stereo. In *ECCV*, 2018. 4
- [40] Botao Ye, Sifei Liu, Haoifei Xu, Li Xueting, Marc Pollefeys, Ming-Hsuan Yang, and Peng Songyou. No pose, no problem: Surprisingly simple 3d gaussian splats from sparse unposed images. In *ICLR*, 2025. 1, 3, 4, 5, 7, 9, 11
- [41] Chandan Yeshwanth, Yueh-Cheng Liu, Matthias Nießner, and Angela Dai. Scannet++: A high-fidelity dataset of 3d indoor scenes. In *ICCV*, 2023. 11
- [42] Zehao Yu, Anpei Chen, Binbin Huang, Torsten Sattler, and Andreas Geiger. Mip-splatting: Alias-free 3d gaussian splatting. *CVPR*, 2024. 4
- [43] Huangying Zhan, Ravi Garg, Chamara Saroj Weerasekera, Kejie Li, Harsh Agarwal, and Ian Reid. Unsupervised learning of monocular depth estimation and visual odometry with deep feature reconstruction. In *CVPR*, 2018. 1, 4
- [44] Shunyuan Zheng, Boyao Zhou, Ruizhi Shao, Boning Liu, Shengping Zhang, Liqiang Nie, and Yebin Liu. Gps-gaussian: Generalizable pixel-wise 3d gaussian splatting for real-time human novel view synthesis. In *CVPR*, 2024. 1, 4
- [45] Tinghui Zhou, Matthew Brown, Noah Snavely, and David G Lowe. Unsupervised learning of depth and ego-motion from video. In *CVPR*, 2017. 1, 4
- [46] Tinghui Zhou, Richard Tucker, John Flynn, Graham Fyffe, and Noah Snavely. Stereo magnification: learning view synthesis using multiplane images. *ACM TOG*, 2018. 3, 7, 8, 9, 11
- [47] Zihan Zhu, Songyou Peng, Viktor Larsson, Weiwei Xu, Hujun Bao, Zhaopeng Cui, Martin R Oswald, and Marc Pollefeys. Nice-slam: Neural implicit scalable encoding for slam. In *CVPR*, 2022. 1, 4
- [48] Zehao Zhu, Zhiwen Fan, Yifan Jiang, and Zhangyang Wang. Fsgs: Real-time few-shot view synthesis using gaussian splatting. In *ECCV*, 2024. 1, 2
- [49] Zihan Zhu, Songyou Peng, Viktor Larsson, Zhaopeng Cui, Martin R Oswald, Andreas Geiger, and Marc Pollefeys. Nicer-slam: Neural implicit scene encoding for rgb slam. In *3DV*, 2024. 1, 4